

# A Taste of Data Science and Machine Learning: a Hands-on Approach



---

Elie Kawerk

# Outline

- Intro to Data Science and Machine Learning.
  - Optimizing functions with gradient descent.
  - Linear regression and Logistic Regression.
  - Overfitting and ways to combat it.
  - Decision Trees and Ensemble Methods.
  - Cross Validation for Model Selection and Hyperparameter Tuning.
-

# Data Science

Data Science is an interdisciplinary field that aims at:

- solving problems by extracting insights from structured and unstructured data,
  - building predictive and descriptive models of data.
-

# Applications of Data Science

- Spam filtering (email)
- Recommender systems (Amazon, Facebook, ...)
- Stock market prediction
- Churn prediction
- Computer vision
- Speech recognition
- Fraud detection
- and more .....

# The Data Science Process

Ask a **question** about your business and how collecting data can help; Set a **target** and **goals**



**Design an experiment** to collect data



**Collect raw data** and **complement** it with other sources such as the web or a database



**Explore, Manipulate, Visualize, Wrangle** and **Clean Data**; **Engineer** new meaningful **Features**



**Build Machine Learning Models** & **Perform Statistical Analyses**



**Communicate** your results clearly through **storytelling** and **visualization**; **Explain** how your analysis helps solving the problem and achieves the set goals

---

# Machine Learning

Machine learning is the subfield of [computer science](#) that, according to [Arthur Samuel](#) in 1959, gives "computers the ability to learn without being explicitly programmed."

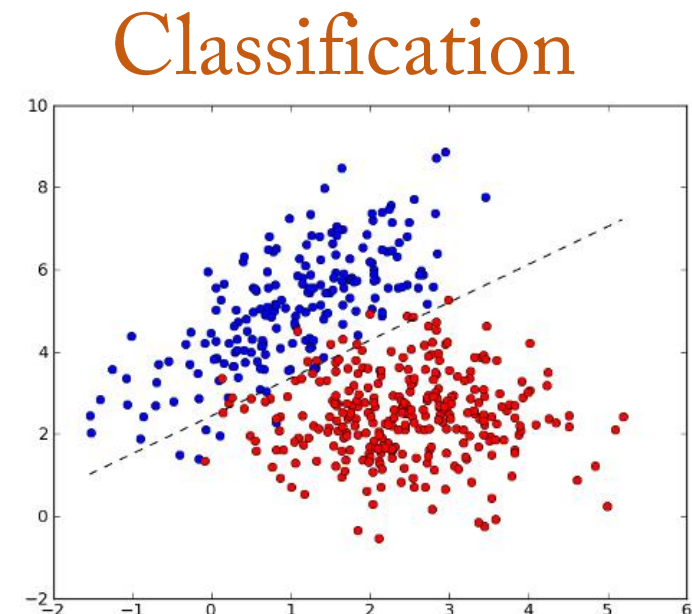
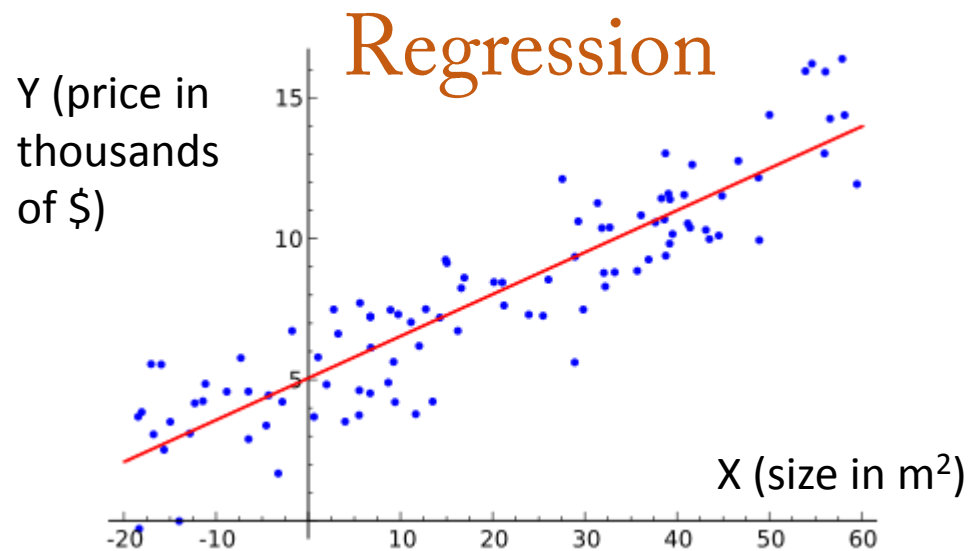
- Wikipedia

---

# Types of Machine Learning

**Supervised learning:** features ( $X$ ) and labels ( $y$ ) are given.  
The task is to learn the mapping  $f$  between  $X$  and  $y$ .


$$f: X \rightarrow y$$




# Types of Machine Learning

Unsupervised learning: only features ( $X$ ) are provided.

## Clustering




## Dimensionality reduction






# Types of Machine Learning

## Reinforcement learning




# Machine Learning vs. Classical Programming

## Traditional Software Engineering



# Machine Learning vs. Classical Programming

## Supervised Machine Learning




# Optimization with Gradient Descent

- Machine learning involves learning a model's parameters from data.
  - In supervised learning, a model must fit data points.
  - In a lot of cases, fitting models to data can be reduced to an optimization problem.
  - Usually, we deal with the optimization of convex functions.
-

# Optimization with Gradient Descent

- In the simplest form, optimization consists of minimizing or maximizing a real function.
- The function may be very complex and multivariate.

# Optimization with Gradient Descent



$J(w)$  denotes the function of  $w$  that we want to minimize

$$w_{\text{optimal}} = \operatorname{argmin}_w J(w)$$

# Optimization with Gradient Descent

## Algorithm

1. Initialize  $X$  with some value which can be random.
2. Pick a learning rate  $\alpha$  i.e. a shrinkage factor to descend along the tangent to the curve.

3. Update  $X$  according to:

$$X = X - \alpha \times \frac{\partial J}{\partial X}$$

---

# Optimization with Gradient Descent

## Algorithm

4. Repeat 2 and 3 for a certain number of iterations or ‘epochs’ hopefully converging at the end.

Note that step 3 should be applied on all the components of vector  $\mathbf{X}$  simultaneously.

---





Let's practice!

# Linear and Logistic Regression

- Now that we know how to optimize functions, it's time to put our knowledge into practice!
  - Linear Regression is one of the most simple models used for regression, that is predicting continuous valued targets.
  - Although it has the term 'regression' in it, Logistic Regression is a linear classification model. It is used for predicting the discrete labels of a target.
-

# Linear Regression


Given the Boston housing dataset (UCI machine learning repository).

X (Average Number of Rooms)

Y (Median House Price)

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

# Linear Regression



# Linear Regression

## Model

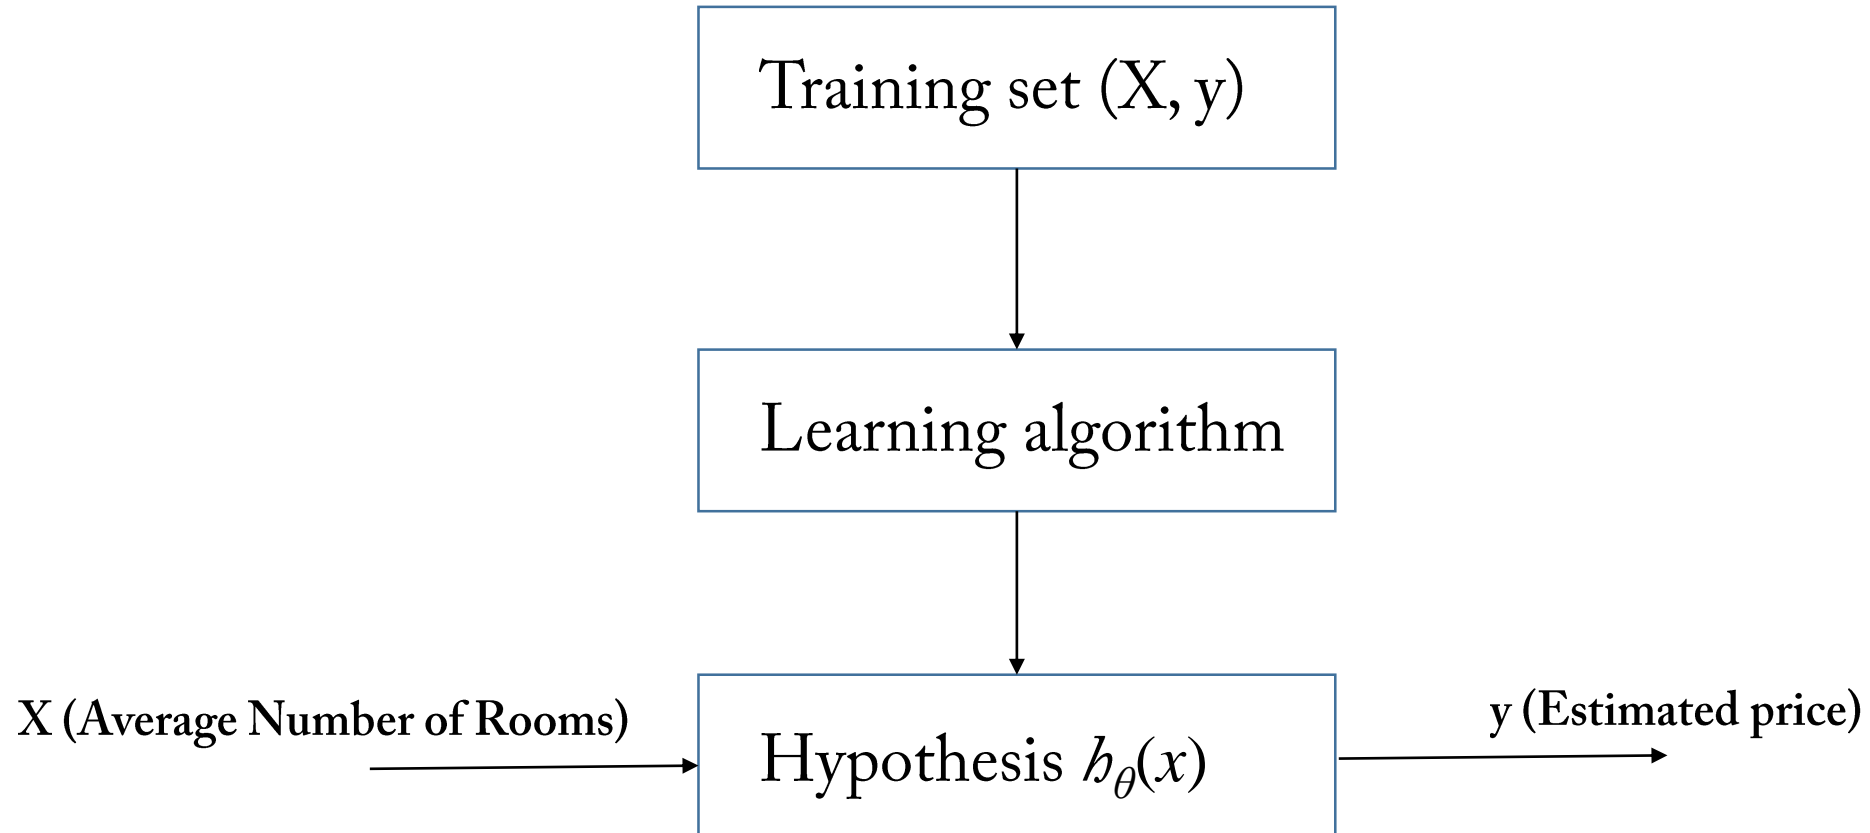
In 1D the linear regression model  $h_{\theta}(x)$  (also known as the hypothesis) has the following form:

$$h_{\theta}(x) = \underbrace{\theta_0}_{\text{intercept}} + \underbrace{\theta_1}_{\text{slope}} \times x$$


The task is to determine the  $\theta_i$  's with  $\theta = (\theta_0, \theta_1)$

---

# Linear Regression



# Cost function



# Cost function for Linear Regression

A cost function measures the agreement between the true labels and the predicted ones. Given a dataset of  $N$  observations.

$$((x_1, y_1), \dots, (x_i, y_i), \dots, (x_N, y_N))$$

The cost function for linear regression is the mean squared error:

$$J(\theta) = \frac{1}{2N} \sum_{i=1}^N (y_i - h_{\theta}(x_i))^2$$



# Gradient Descent for Linear Regression

In order to find the optimal hypothesis that fits the data, we need to minimize the cost function with respect to  $\theta$ .

$$\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

This minimization is performed using gradient descent.

---

# Gradient Descent for Linear Regression

## Algorithm

Repeat until convergence:

$$\theta_j := \theta_j - \underbrace{\alpha}_{\text{learning rate}} \times \underbrace{\frac{\partial J(\theta_j)}{\partial \theta_j}}_{\text{partial derivative along } \theta_j}$$

# Gradient Descent for Linear Regression

For linear regression, we can compute the derivative of the cost function to obtain the update we have to make:

$$\theta_j := \theta_j - \frac{\alpha}{N} \sum_{i=1}^N (h_{\theta}(x_i) - y_i) \times x_i^j$$

with  $j = 0, 1$ ,  $x_i^0 = 1$  and  $x_i^1 = x$

---



Let's practice!

# Logistic Regression

## Model

Logistic Regression aims at predicting the probabilities of labels. For a binary classification problems 0 denotes the negative class and 1 denotes the positive class. The hypothesis is the probability of predicting class 1 given a feature vector  $x$ .

$$h_{\theta}(x) = \frac{1}{1 + e^{-(\theta_0 x_0 + \dots + \theta_D x_D)}}$$

# Cost function for Logistic Regression

Given a dataset of  $N$  observations.

$$((x_1, y_1), \dots, (x_i, y_i), \dots, (x_N, y_N))$$

with  $y_i = 0$  or  $1$

The cost function for logistic regression is given by:

$$J(\theta) = - \sum y_i \log(h_\theta(x_i)) + (1 - y_i) \log(1 - h_\theta(x_i))$$

---



Let's practice!

# Overfitting

- Overfitting is fitting the training data more than is warranted.
  - When you overfit, your model performs well on the training data but fails to generalize on unseen data.
  - A model is said to suffer from **high bias** if the hypothesis is **not complex enough** to capture the pattern of the signal in the data.
  - On the other hand a **high variance** model is so complex that, not only it fits the signal, but also fits the **noise** present in the training set!
-



# Overfitting






Figure taken from Python Machine Learning, Sebastian Raschka, Packt 2015

# Overfitting



# Combating Overfitting

- One way to combat overfitting is to add a regularization term.
- Regularization not only reduces overfitting but also handles collinearity and filters out noise from data.
- In  $L2$  regularization, we add the following term to the cost function:

$$\frac{\lambda}{2} ||\theta||^2 = \frac{\lambda}{2} \sum_{j=1}^N \theta_j^2$$

# Combatting Overfitting

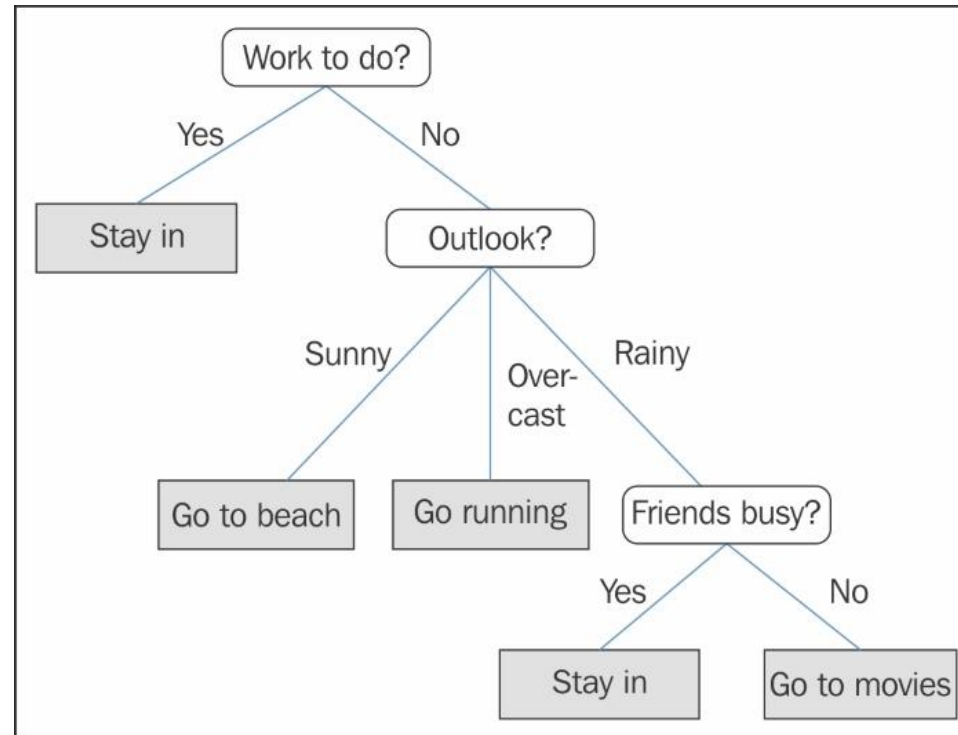
- Regularization has the effect of shrinking the weights of some features.



Let's practice!

# Decision Trees

Make decision based on a series of questions.



# Decision Trees

Decision Trees learn by maximizing information gain. At each split, a feature and a split point are selected such that they maximize:

$$IG(D_p, f) = I(D_p) - \sum_{j=1}^m \frac{N_j}{N_p} I(D_j)$$

Using the entropy as a measure of information in a node:


$$I_H(t) = - \sum_{i=1}^c p(i|t) \log_2 p(i|t)$$

# Decision Trees

- Decision trees can build complex decision boundaries by dividing the feature space into rectangles.
- However, we have to be careful since the deeper the decision tree, the more complex the decision boundary becomes, which can easily result in overfitting.
- To reduce overfitting, we should **prune** the tree. That is we impose a maximal depth on its growth.



# Decision Trees



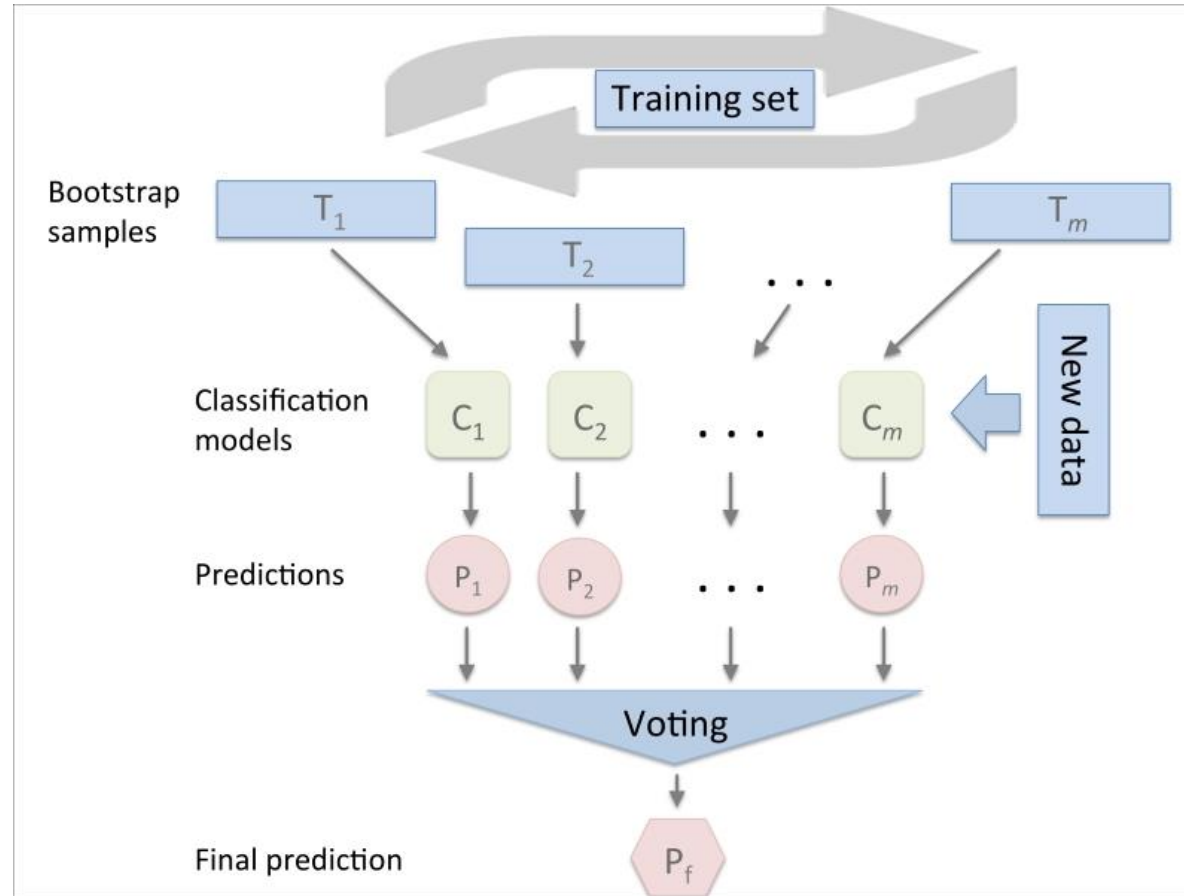


Let's practice!

# Ensemble Methods

- In ensemble methods, we combine the predictions of different classifiers via a meta-classifier.
  - Think of it this way: an individual expert may give the wrong prediction.
  - On the other hand, when the judgements of many experts are combined strategically, the prediction is more robust and less prone to error.
-

# Ensemble Methods: Bagging



# Ensemble Methods: Random Forests

## Algorithm

1. Draw a random bootstrap sample of size  $n$ .
  2. Grow a decision tree from the sample. At each node:
    - a. Randomly select  $d$  features without replacement.
    - b. Split the node using the feature that provides the best split.
  3. Repeat 1 to  $k$  times.
  4. Aggregate results.
-



Let's practice!

# Evaluating Classification Models

- Consider a binary classification problem where there is a class imbalance; that is the negative class has a much higher number of instances than the positive class. Ex: 90 % negative.
- If your evaluation metric is accuracy, then predicting the negative class blindly would achieve a score of 90 %!
- There are many methods used to deal with class imbalance. We will discuss a method in which we use evaluation metrics other than accuracy.

# Evaluating Classification Models

		Predicted class	
		<i>P</i>	<i>N</i>
Actual Class	<i>P</i>	True Positives (TP)	False Negatives (FN)
	<i>N</i>	False Positives (FP)	True Negatives (TN)

$$ERR = \frac{FP + FN}{FP + FN + TP + TN}$$

$$ACC = \frac{TP + TN}{FP + FN + TP + TN} = 1 - ERR$$

$$FPR = \frac{FP}{N} = \frac{FP}{FP + TN}$$

$$TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$



# Evaluating Classification Models

- Precision is a metric that measures the rate of true positive among all predicted positives:

$$PRE = \frac{TP}{TP + FP}$$

- Recall is another name of the TPR:  $REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$

- In practice, we use the f1-score defined as follows:

$$F1 = 2 \frac{PRE \times REC}{PRE + REC}$$

# Cross-validation

- Evaluating your model on the training set gives a biased estimate of your model's skill.
  - This is because you are evaluating the model on data that it has already seen.
  - Cross-validation is a technique used to obtain a less biased estimate of the skill achieved by your model.
-

# K-fold Cross-validation

- For small to medium sized datasets, the gold standard is K-fold cross validation.
  - Split the training data into K folds.
  - Train the model K times to obtain a list of K scores.
  - At each iteration, train the model on (K-1) folds and evaluate it on an individual fold unseen by in training.
-

# K-fold Cross-validation



The average score of a model obtained using K-fold CV estimates its generalization score.

# Model Selection


- After training many models using K-fold cross validation, the winning model should be declared as the one that achieves the highest mean CV score.
- Note that to report a final **unbiased** estimate of your model's skill, you should **evaluate** the model on a **test set** that it has never seen. This step is **crucial**!
- Do not compromise the last step! otherwise you are **snooping** the data.

# Hyperparameter Tuning

- The parameters of your model that are not learned through training are called hyperparameters.
- A model's hyperparameters should be tuned so that the model achieves the best results.
- Examples of hyperparameters include: the regularization term in linear regression, tree depth and number of estimators in random forests.

# Hyperparameter Tuning

- Hyperparameter tuning is performed through grid search.
- Think of it as adjusting the control knobs of a thermostat to obtain the 'best' temperature.



---

Thank you for your attention!

&

Let's practice!

---